

# DATABASE


---

Martin Brugnara

$2\{0,1\}/11/2019$

[mb@disi.unitn.eu](mailto:mb@disi.unitn.eu)





[wiki.postgresql.org/wiki/Logo](https://wiki.postgresql.org/wiki/Logo)

# TODAY

1. Interacting with a DMBS
2. Non trivial features of PG
3. Database course Q&A

## Command

```
psql -h <server> -U <user> <database>
```

## Remote server

```
psql -h sci-didattica.unitn.it -U db_001 db_001
```

## Local server

```
psql db_001
```

# INSTALL

## Client only

- Debian/Ubuntu: `postgresql-client`
- Brew: `libpq`

## Server & Client

- Debian: `postgresql`
- Brew: `postgresql`
- macOS: `Potgres.app`

## GUI

- web app: `pgAdmin 4`

The server is reachable only from the local network  
`unitn`, `unitn-x`, (local) eduroam, cable

[wiki.unitn.it/pub:conf-vpn](https://wiki.unitn.it/pub:conf-vpn)

# ACCOUNTS

gspreadsheets: 1627VkH7F01MueeWfsJ5j9V\_gqCDs6SSxYeEyCmbzRf0

## Structure

- Database: db\_ddd
- Username: db\_ddd
- Password: secret\_ddd

Where ddd in [001, 250].

## Change your password

*# Connect via psql*

```
PGPASSWORD=secret_001 psql \
```

```
    -h sci-didattica.unitn.it -U db_001 db_001
```

*# Change password via query*

```
ALTER USER "db_001" WITH PASSWORD 'newpass';
```



### Quotation marks

" " : names of schemas, tables, columns, *etc.*

' ' : string constant.

### Cast

— *SQL standard*

**CAST** ( expression **AS** type )

— *Old style PG*

expression :: type

### Auto increment

SERIAL – see documentation.

### **Case Sensitivity**

Tables, columns names: yes with ""

...without, transformed to lowercase.

String: case sensitive.

### **Naming things** - doc.

Database Cluster (Server - User & Groups)

→ Databases (Name Space) → Schemas → Tables.

## JDBC & psycopg2

---

# JDBC - CONNECT

```
/* JDBCExample.java */  
import java.sql.*;  
  
/* class JDBCExample > static void main () */  
String url = "jdbc:postgresql://sci-didattica.unitn.it/db_001";  
Properties props = new Properties();  
props.setProperty("user", "db_001");  
props.setProperty("password", "secret_001");  
props.setProperty("ssl", "false");  
Connection conn = DriverManager.getConnection(url, props);  
/* Do stuff. */  
conn.close();  
  
# Download jdbc.postgresql.org/download/postgresql-42.2.8.jar  
java -cp "$(pwd)/postgresql-42.2.8.jar;$(pwd)/" JDBCExample
```

<https://jdbc.postgresql.org/documentation/head>

# JDBC - QUERY

```
Statement st = conn.createStatement();
ResultSet rs = st.executeQuery("SELECT_*_FROM_mytable_WHERE_column");
while (rs.next())
{
    System.out.print("Column_1_returned_");
    System.out.println(rs.getString(1));
}
rs.close();
st.close();
```

<https://jdbc.postgresql.org/documentation/head/query.html>

# JDBC - PREPARED STATEMENT

```
String q = "INSERT INTO tbl_(foo) VALUES(?)";  
PreparedStatement st = conn.prepareStatement(q);  
st.setObject(1, new_value);  
st.executeUpdate();  
st.close();
```

[jdbc.postgresql.org/documentation/head/update.html](https://jdbc.postgresql.org/documentation/head/update.html)  
[jdbc.postgresql.org/documentation/head/java8-date-time.html](https://jdbc.postgresql.org/documentation/head/java8-date-time.html)

# PYTHON

```
import psycopg2                                # initd.org/psycopg/

# Connect
conn = psycopg2.connect("dbname='db_001' _user='db_001' _" +
    " host='sci-didattica.unitn.it' _password='secret_001'")
cur = conn.cursor()

# Create / Prepared statement
cur.execute("CREATE _TABLE _test _...")
cur.execute("INSERT _INTO _test _ (num, _data) _VALUES _(%s, _%s)",
    (100, " abc ' def" ))
conn.commit()                                # Make changes permanent

# Query
cur.execute("SELECT _*_ _FROM _test ;")
cur.fetchone()                                # next row only
cur.fetchall()                                # all remaining ones

# Clean up
cur.close()
conn.close()
```

## Nice types

---



# RANGE

## Types

numrange, tstzrange, daterange

## Syntax

```
'(7,12]' :: numrange
```

```
numrange(7, 12, '()')
```

# RANGE

## Types

numrange, tstzrange, daterange

## Syntax

```
'(7,12]' :: numrange  
numrange(7, 12, '()')
```

## Operators

[www.postgresql.org/docs/current/functions-range.html#](http://www.postgresql.org/docs/current/functions-range.html#RANGE-OPERATORS-TABLE)

RANGE-OPERATORS-TABLE

@>, &&, >> &>, -|- , +, \*, -

# RANGE

## Types

numrange, tstzrange, daterange

## Syntax

```
'(7,12]' :: numrange  
numrange(7, 12, '()')
```

## Operators

[www.postgresql.org/docs/current/functions-range.html#](http://www.postgresql.org/docs/current/functions-range.html#RANGE-OPERATORS-TABLE)

RANGE-OPERATORS-TABLE

@>, &&, >> &>, -|- , +, \*, -

## Functions

[www.postgresql.org/docs/current/functions-range.html#](http://www.postgresql.org/docs/current/functions-range.html#RANGE-FUNCTIONS-TABLE)

RANGE-FUNCTIONS-TABLE

lower, upper, isempty, lower\_inc, lower\_inf, range\_merge.

# ARRAYS

## Declaration

```
field1    type [ ] ,  
field2    type [ ] [ ] ,  
field3    type [ n ] [ m ] [ k ] ,  
... 
```

## Syntax

```
' { 1 , 2 , 3 , 4 } '  
' { { 1 , 2 , 3 } , { 4 , 5 , 6 } , { 7 , 8 , 9 } } '  
ARRAY [ 1 , 2 , 3 , 4 ]
```

## Navigation

```
field1 [ 0 ] , field2 [ 2 : 4 ] [ : 3 ]
```

# ARRAYS

## Editing

```
field1[0] = 33  
field2[1][:3] = '{2,3,4}'  
arr_1 || arr_2  
arr_1 || val || NULL  
array_cat()
```

## Search

```
ANY, ALL, &&  
array_position , array_positions
```

[www.postgresql.org/docs/current/arrays.html](http://www.postgresql.org/docs/current/arrays.html)

## JSON & JSONb

### Navigation

```
->      # descend one node, returns JSON node  
->>     # descend one node, returns value/JSON as text  
#>     # query path, returns JSON node  
#>>    # query path, returns value/JSON as text
```

## JSON & JSONb

**Filter / Modify** JSONb only.

Containment: <@, ?, ?|

Concat: ||

**Delete:** -, #-

Path query (& compare): @?, @@

Text only, dictionary.

Mostly replaced by JSONb.



## Cool operators and functions

---

# WINDOW FUNCTIONS

Like aggregation, but does not collapse rows.

```
func() [FILTER ( WHERE filter_clause )] OVER (window_def)
```

## Additional functions <sup>1</sup>

[www.postgresql.org/docs/12/functions-window.html](http://www.postgresql.org/docs/12/functions-window.html)

- row\_number()
- lag, lead (expression, offset, default)
- first\_value, last\_value, (expression)
- nth\_value (expression, nth)

# WINDOW FUNCTIONS

Like aggregation, but does not collapse rows.

```
func() [FILTER ( WHERE filter_clause )] OVER (window_def)
```

## Window definition <sup>2</sup>

- PARTITION BY
- ORDER BY
- frame\_clause\*

[www.postgresql.org/docs/12/sql-expressions.html#](http://www.postgresql.org/docs/12/sql-expressions.html#SYNTAX-WINDOW-FUNCTIONS)

SYNTAX-WINDOW-FUNCTIONS

# WINDOW FUNCTIONS

Like aggregation, but does not collapse rows.

```
func() [FILTER ( WHERE filter_clause )] OVER (window_def)
```

## Usage examples

“First  $n$  items for each group”, “running sums”, “stats”, ...

Buy at least 4 and the 2 less expensive are free.

```
SELECT customer_id , SUM(price) FROM (  
  SELECT customer_id , price ,  
    ROW_NUMBER() as r  
    OVER (PARTITION BY customer_id ORDER BY price),  
    COUNT() as m OVER (PARTITION BY customer_id),  
) tmp WHERE m < 4 OR (m >= 4 AND r > 2)
```

## Series Generating Functions

`generate_series(start , stop , step)`

`generate_subscripts(array , dim , reverse)`

[www.postgresql.org/docs/current/functions-srf.html](http://www.postgresql.org/docs/current/functions-srf.html)

# Indexes

---

# INDEXES OVERVIEW

- B-tree (default)
- Hash
- GiST
- SP-GiST
- GIN
- BRIN

[www.postgresql.org/docs/current/indexes-types.html](http://www.postgresql.org/docs/current/indexes-types.html)

## Speeds up

- `<, <=, =, >=, >`  
→ BETWEEN, IN
- LIKE  
constant pattern anchored to the beginning, e.g. 'foo%'
- IS NULL, IS NOT NULL



**Speeds up**

- =

# GENERALIZED SEARCH TREE (GIST)

Balanced, tree-structured access method.

Template in which to implement arbitrary indexing schemes.

## **Natively supports**

- Geometry: box, circle, point, polygon
- Internet Addresses: inet, cidr
- any range type
- tsquery, tsvector

document representation to optimized full text search

Many more from the contrib collection.

[www.postgresql.org/docs/current/gist-intro.html](http://www.postgresql.org/docs/current/gist-intro.html)

Similar to GiST but support non-balanced disk-based data structures.

*e.g.* quadtrees, k-d trees, and radix trees (tries)

Inverted indexes used for complex objects.

## **Natively supports**

- array
- jsonb
- tsvector

[www.postgresql.org/docs/current/gin-builtin-opclasses.html](http://www.postgresql.org/docs/current/gin-builtin-opclasses.html)

# BLOCK RANGE INDEXES (BRIN)

Used for very large tables.

Indexes range of consecutive table's physical blocks.

[www.postgresql.org/docs/current/brin-intro.html](http://www.postgresql.org/docs/current/brin-intro.html)

**Explain!**

---

# THE EXPLAIN COMMAND

“show the execution plan of a statement”

For each query node shows

e.g. (cost=0.00..0.24 rows=10 width=82)

- cost=f..a

f: Cost up to the first row

a: Cost for computing all rows

- rows

Expected number of returned rows

- width

Expected width of the rows

[www.postgresql.org/docs/current/sql-explain.html](http://www.postgresql.org/docs/current/sql-explain.html)

# EXPLAIN ANALYZE

Executes the command and  
shows actual run times and other statistics.

For each query node shows

e.g. ... (actual time=0.025..0.029 rows=10 loops=1)

- actual time  
Average execution time for one loop
- rows  
Actual number of returned rows
- loops  
Times the node has been executed



# EXPLAIN ANALYZE EXAMPLE 1

```
EXPLAIN ANALYZE  
SELECT surname , name  
FROM people_person LIMIT 10;
```

# EXPLAIN ANALYZE EXAMPLE 1

```
EXPLAIN ANALYZE
```

```
SELECT surname, name
```

```
FROM people_person LIMIT 10;
```

```
Limit (cost=0.00..0.24 rows=10 width=15)
```

```
(actual time=0.016..0.022 rows=10 loops=1)
```

```
-> Seq Scan on people_person
```

```
(cost=0.00..588.88 rows=24988 width=15)
```

```
(actual time=0.014..0.017 rows=10 loops=1)
```

```
Planning time: 0.141 ms
```

```
Execution time: 0.043 ms
```

## EXPLAIN ANALYZE EXAMPLE 2 - QUERY

```
EXPLAIN ANALYZE
SELECT p.surname, p.name, c.name,
       EXTRACT(YEAR FROM AGE(birth_date)),
FROM people_person p JOIN people_country c
     ON citizenship_id=c.id
WHERE EXTRACT(YEAR FROM AGE(birth_date)) > 25;
```

## EXPLAIN ANALYZE EXAMPLE 2 - RESULT

```
Hash Join  (cost=8.46..1098.43 rows=8329 width=34)
  (actual time=0.294..69.545 rows=14518 loops=1)
Hash Cond: (p.citizenship_id = c.id)
-> Seq Scan on people_person p
  (cost=0.00..963.70 rows=8329 width=23)
  (actual time=0.019..39.535 rows=14596 loops=1)
    Filter: (date_part('year'::text, age((CURRENT_DATE)::timestamp, p.birthdate)) < 18)
    Rows Removed by Filter: 10392
-> Hash  (cost=4.87..4.87 rows=287 width=15)
  (actual time=0.250..0.251 rows=287 loops=1)
  Buckets: 1024  Batches: 1  Memory Usage: 22kB
    -> Seq Scan on people_country c
      (cost=0.00..4.87 rows=287 width=15)
      (actual time=0.008..0.125 rows=287 loops=1)

Planning time: 0.550 ms
Execution time: 70.285 ms
```

**Never ending**

---

- Foreign data wrappers

`www.postgresql.org/docs/current/ddl-foreign-data.html`

- PostGIS

`postgis.net`

- Postgres Weekly

`postgresweekly.com`

- ...

## Q&A

---